

PeachyDB: a miniature relational database

- Author: Fang Han
- 2019.11
- Database Systems @ NYU Courant

TABLE OF CONTENTS

1. [SPECIAL INSTRUCTIONS FOR GRADERS](#)
2. [LIST OF QUERIES](#)
3. [SETUP](#)
 - i. [compile and run with maven](#)
 - ii. [run with shell script](#)
4. [DOCUMENTATION](#)
 - i. [table naming convention](#)
 - ii. [I/O queries](#)
 - iii. [algebraic queries](#)
 - iv. [aggregate queries](#)
 - v. [moving aggregate queries](#)
 - vi. [index](#)
 - vii. [utility queries](#)
5. [FEATURES](#)
6. [STATISTICS](#)

SPECIAL INSTRUCTIONS FOR GRADERS

- step 0: unzip the `.rpz` file
- step 1: put all **input data** files under `input/`
- step 2: put the file containin test queries under `input/`
- step 3: open `input/input_pipe` , change the 2nd line to match the name of the query file above
- step 4: at root dir, run `./run.sh`
- After the above steps, find outputs under `output/`

QUERIES SUPPORTED

1. [showtables](#)
2. [showschema](#)
3. [quit](#)
4. [inputfromfile](#)
5. [outputtofile](#)
6. [select](#)
7. [project](#)
8. [join](#)
9. [concat](#)
10. [sort](#)
11. [count](#)
12. [sum](#)
13. [avg](#)
14. [countgroup](#)
15. [sumgroup](#)
16. [avggroup](#)
17. [movavg](#)
18. [movsum](#)
19. [hash](#)
20. [btree](#)

SETUP

compile and run with maven

- download repo

```
$ git clone https://github.com/TakaiKinoko/PeachyDB.git
```

- compile

```
$ cd PeachyDB
$ mvn compile
```

- build jar

```
$ mvn package
```

- run interactively

```
$ java -cp target/peachyDB-1.0.jar Entry
```

- exit the database
type **quit** when the database is running.

run with shell script

- **./run.sh** at root
- this will feed all query lines from **input/handout** to the database and direct **stdout** to **output/fh643_All0perations**

DOCUMENTATION

table naming convention

- has to start with an alphabetic letter
- syntax using regular expression: **([a-zA-Z]+(.)*)**
- derivative tables:
 - **definition:** tables that are built 'on top of' another (more than one) existing table
 - to differentiate the derivative table columns from its parent(s), it's column names have the format of **<table_name>_<column_name>**
 - queries on the derivative tables should make sure that the columns are addressed according to the rule above

I/O

read from file

- syntax: **<table_name> := inputfromfile(<filepath>)**
- implementation: under **src/io/I0.java**
- note:
 - i. a **<filepath>** must be assigned to a **<table_name>**
 - ii. the database at default tries to read files from the **/input** folder. So **<filepath>** should be the relative path from **/input** to the file
 - iii. reading in a new file will create a new table.
 - iv. a **truncated view** of the table will be printed out to StdOut once data has been read in successfully, for example:

reading from file: input/sales2.txt into table: S...

saleid	I	C	S	T	Q	P
3506	13517	16566	45	73	19	expensive
78345	10528	4745	20	73	23	supercheap
79991	6715	707	75	41	34	expensive
90466	6697	8397	83	92	16	outrageous
22332	9639	2435	29	17	31	moderate
95047	11877	2020	44	79	29	supercheap
48867	12387	15274	98	76	35	supercheap
22220	10650	5746	57	73	24	outrageous
53696	9958	11849	85	16	9	supercheap
34328	11376	4042	50	66	44	supercheap
...

62617	10689	15710	3	73	29	supercheap
74088	6099	14086	37	95	44	moderate
66449	10137	2465	41	73	31	cheap
11662	9096	19072	6	16	21	supercheap
33022	6259	5746	54	11	44	supercheap
86141	10713	5746	71	73	4	outrageous
64366	8775	18198	43	61	49	supercheap
41918	10898	18816	61	92	18	moderate
43539	8229	16589	14	92	47	supercheap
2356	8909	14012	32	82	24	supercheap

Number of entries: 100000

Time cost: 0.1450 seconds

- example: `inputfromfile(sales1.txt)` , where `sales1.txt` is stored inside `/input`

write table to file

- syntax: `outputtofile(<table>, <filename>)`
- implementation: under `src/io/I0.java`
- note:
 - the database at default tries to save files to the `/output` folder.
 - PrettyPrinter** (see `/src/util/PrettyPrinter.java`) is used to format the output table.
 - sample pretty-printed result:

groupby_pricerange	avg_qty
cheap	20.546875
expensive	24.954545454545453
moderate	22.384615384615383
outrageous	23.717047451669597
supercheap	26.10126582278481

Number of entries: 5

algebraic

select

- syntax: `<target_table> := select(<from_table>, <condition1> [and/or <condition2>])`
- the `[and/or <condition2>]` part is optional, which means this select operation takes one or two conditions
- svntax of the condition: `(Column | Constant) [+|-|*|/ Constant] (< | <= | > | >= | != | =) (Column | Constant) [+|-|*|/ Constant]`
- within each condition, the `[+|-|*|/ Constant]` part is optional
- implemented in `src/algebra/Select.java`
- entries selected will be deep copy from the source table
- if a column within the conditions is indexed upon (by either Hash or BTree), the index will be used to perform selection

project

- syntax: `<target_table> := project(<from_table>, <col1>, ..., <coln>)`
- implemented in `src/algebra/Project.java`
- actually fulfilled by the function `projectTable` in `src/db/Database.java`
- columns selected will be **shallow copy** (pointer) of the source table

join

- syntax: `<target_table> := join(<table1>, <table2>, <condition1> [and/or <condition2>])`
- the `[and/or <condition2>]` part is optional, which means this join operation takes one or two conditions
- svntax of the condition: `<table_name1>.<column_name1> ([+|-|*|/] <constant1>) [>|<|!=|=|>=|<=] <table_name2>.<column_name2> [+|-|*|/] <constant2>`

- within each condition, the `([+|-|*|/] <constant>)` part is optional

- implemented in `src/algebra/Join.java`

concat

- syntax: `<target_table> := concat(<table1>, <table2>)`
- implemented in `src/algebra/Concat.java`
- actually fulfilled by the function `concatTables` in `src/db/Database.java`

sort

- syntax: `<target_table> := sort(<from_table>, <col1>, ..., <coln>)`
- implemented in `src/util/Sort.java`

aggregate

count

- syntax: `<to_table> := count(<from_table>, <column_name>)`
- implemented in `src/aggregation/Aggregate.java`

sum

- syntax: `<to_table> := sum(<from_table>, <column_name>)`
- implemented in `src/aggregation/Aggregate.java`

avg

- syntax: `<to_table> := avg(<from_table>, <column_name>)`
- implemented in `src/aggregation/Aggregate.java`

countgroup

- count the number of entries of a column from a table grouped on an ordered list of columns serving as grouping conditions
- syntax: `<to_table> := countgroup(<from_table>, <column_name>, <groupby_col1>, ..., <groupby_coln>)`
- implemented in `src/aggregation/GroupAgg.java`
- based on internal method `groupby` implemented in `src/aggregation/GroupAgg.java`

sumgroup

- compute the sum of a column from a table grouped on an ordered list of columns serving as grouping conditions
- syntax: `<to_table> := sumgroup(<from_table>, <column_name>, <groupby_col1>, ..., <groupby_coln>)`
- implemented in `src/aggregation/GroupAgg.java`
- based on internal method `groupby` implemented in `src/aggregation/GroupAgg.java`

avggroup

- compute the average of a column from a table grouped on an ordered list of columns serving as grouping conditions
- syntax: `<to_table> := avggroup(<from_table>, <column_name>, <groupby_col1>, ..., <groupby_coln>)`
- implemented in `src/aggregation/GroupAgg.java`
- based on internal method `groupby` implemented in `src/aggregation/GroupAgg.java`

moving aggregates

moving average

- syntax: `<toTable> := movavg(<fromTable>, <col>, <window_len>)`
- implemented in `src/aggregation/Moving.java`
- fulfilled by private internal method `apply` within `src/aggregation/Moving.java`

moving sum

- syntax: `<toTable> := movsum(<fromTable>, <col>, <window_len>)`
- implemented in `src/aggregation/Moving.java`
- fulfilled by private internal method `apply` within `src/aggregation/Moving.java`

index

hash

- syntax: `Hash(<table>, <column>)`
- implemented in `src/index/Hash.java` through Java's native `HashMap` class

btree

- syntax: `Btree(<table>, <column>)`
- implemented in `src/index/Btree.java`
- Btree implementation: `src/btree`

utility

quit

- syntax: `quit` or `Quit`
- implemented in `src/io/QueryParser.java`

show tables

- syntax: `showtables()`
- implemented in `src/db/Database.java`
- sample output:

Table	Size
R2	900
R	1000
S	100000
T	3642
T2prime	391
T1	391
T2	391
R1	900
T3	391

show schemas

- syntax: `showschema()`
- implemented in `src/db/Database.java`
- sample output:

Table	Schema
R	saleid itemid customerid storeid time qty pricerange
S	saleid I C S T Q P

FEATURES

Pretty-Printer

- implemented in `src/util/PrettyPrinter.java`

STATISTICS

- line counts using: `$ find . -name '*.java' | xargs wc -l`

```
232 ./src/main/java/aggregation/GroupAgg.java
151 ./src/main/java/aggregation/Moving.java
 85 ./src/main/java/aggregation/Aggregate.java
 61 ./src/main/java/util/Sort.java
218 ./src/main/java/util/PrettyPrinter.java
102 ./src/main/java/util/GroupKey.java
159 ./src/main/java/util/Cond.java
142 ./src/main/java/util/Utls.java
 35 ./src/main/java/util/SortGroupKeyMap.java
185 ./src/main/java/io/IO.java
```

```
218 ./src/main/java/io/QueryParser.java
186 ./src/main/java/parser/Parser.java
18  ./src/main/java/btree/BTKeyValue.java
1019 ./src/main/java/btree/BTree.java
11  ./src/main/java/btree/BTIteratorIF.java
26  ./src/main/java/btree/BTException.java
61  ./src/main/java/btree/BTNode.java
42  ./src/main/java/btree/SimpleFileWriter.java
58  ./src/main/java/db/DynamicTable.java
187 ./src/main/java/db/Table.java
380 ./src/main/java/db/Database.java
39  ./src/main/java/index/BTTestIteratorImpl.java
84  ./src/main/java/index/Btree.java
65  ./src/main/java/index/Hash.java
65  ./src/main/java/index/BtreeKey.java
71  ./src/main/java/Entry.java
354 ./src/main/java/algebra/Join.java
587 ./src/main/java/algebra/Select.java
55  ./src/main/java/algebra/Project.java
34  ./src/main/java/algebra/Concat.java
4930 total
```